

MATLAB for High Performance Computing

Shao Hao Chen

Research Computing Services

Information Services and Technology

Boston University

Why using Matlab for HPC?

- ❖ Many Matlab programs run faster on high-performance computing (HPC) clusters than on regular laptops/desktops.
- ❖ Boston University (BU) Shared Computing Cluster (SCC) is an HPC cluster with over 11,000 CPU cores and over 250 GPUs.
- ❖ MATLAB site license is available to all BU users. (Unlimited number on SCC).

- ❖ Matlab programs can be exceptionally fast if they are well-designed, and painfully slow if not. It is necessary to optimize MATLAB codes to obtain good performance.
- ❖ The code-optimized methods to be mentioned in this tutorial can be used not only on HPC clusters but also on regular laptops/desktops.

- ❖ The skills you learn today should enable you to solve bigger problems faster using MATLAB.

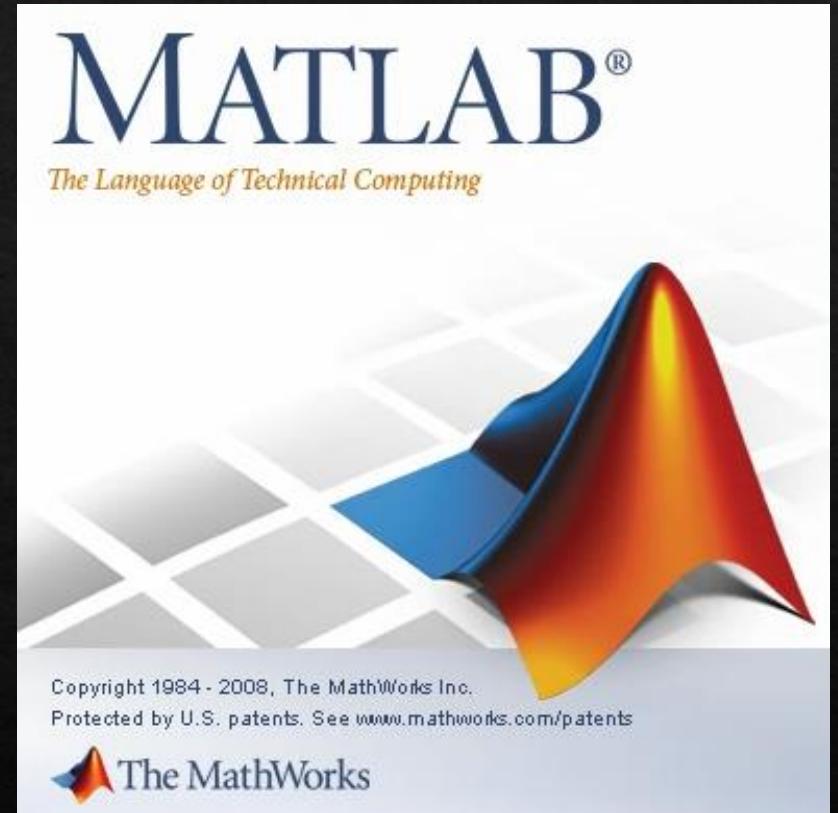
Outline

- ❖ Use Matlab on BU SCC

- ✓ Start up
- ✓ Submit Matlab jobs
- ✓ Distributed jobs

- ❖ Optimize Matlab codes

- ✓ Remove unnecessary works
- ✓ Optimize memory usage
- ✓ Use optimized built-in functions/operators



Access to BU SCC resources

- ❖ Log in to SCC:

```
$ ssh -X username@sccl.bu.edu
```

- ❖ Interactive session: working interactively on compute nodes.

```
$ qrsh      # Start an interactive session
```

- ✓ Any job costing more than 15-minute CPU time on the login node will be killed.

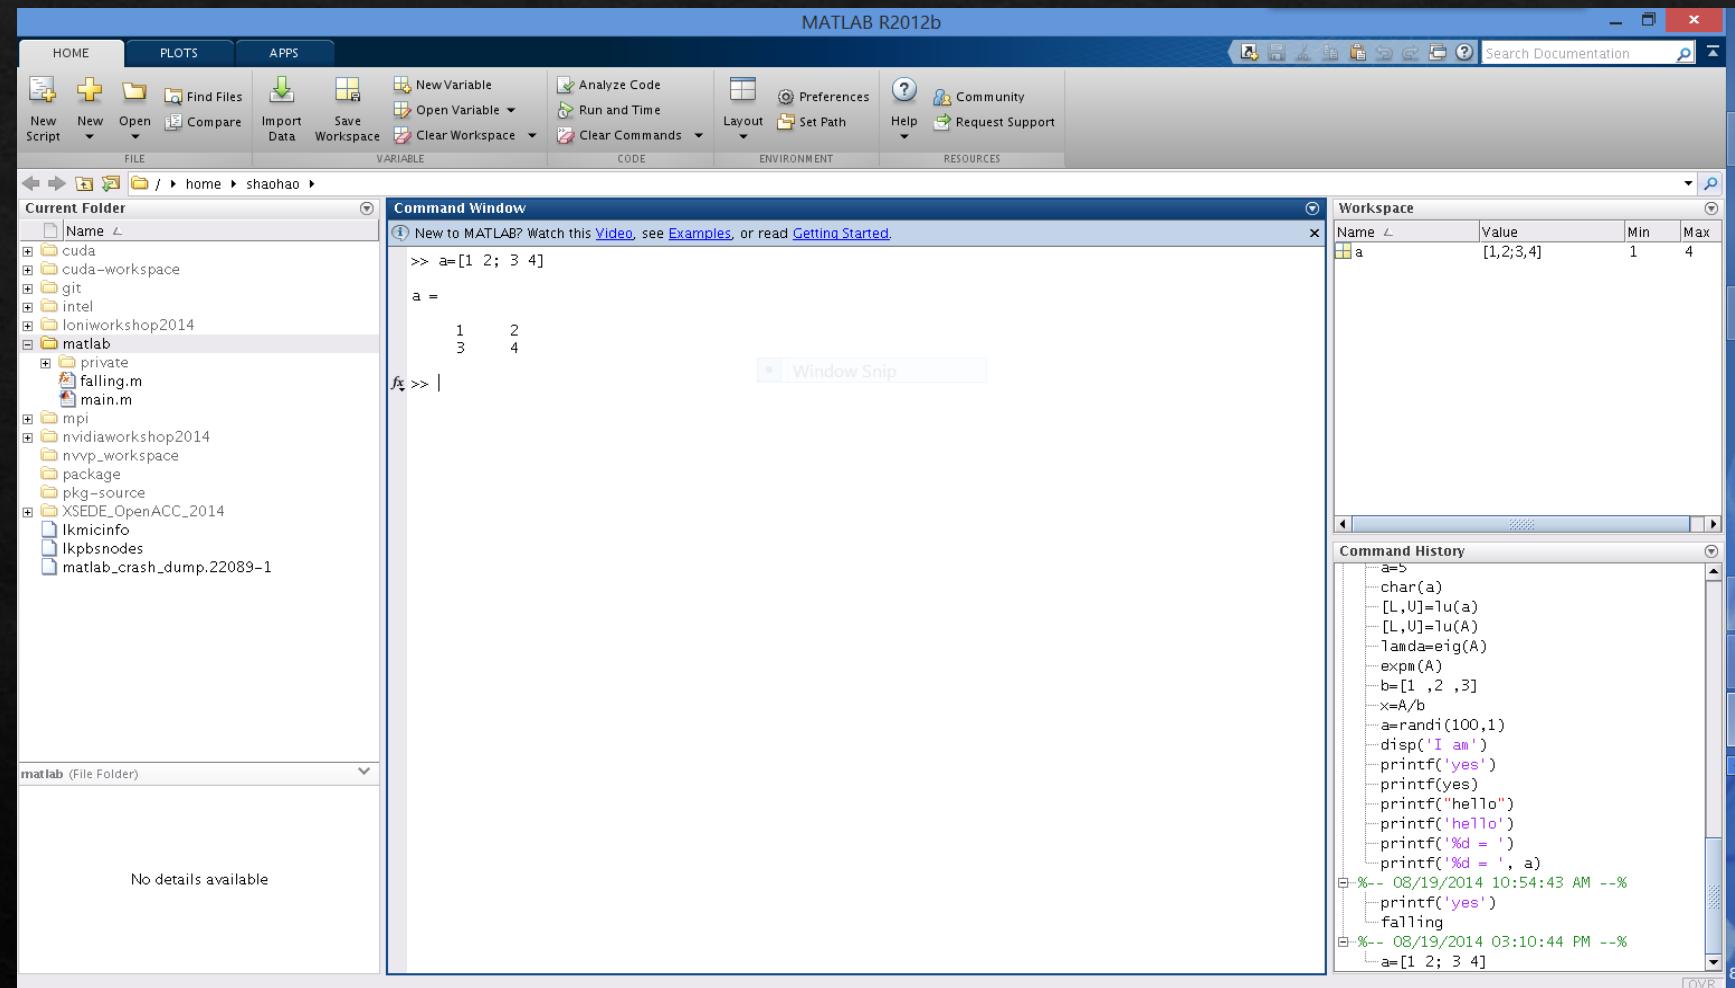
- ❖ Load Matlab module:

```
$ moduleavail | grep matlab    # See all available versions
```

```
$ module load matlab            # Set up environment variables
```

Graphic platform

- Open Matlab
- \$ matlab &

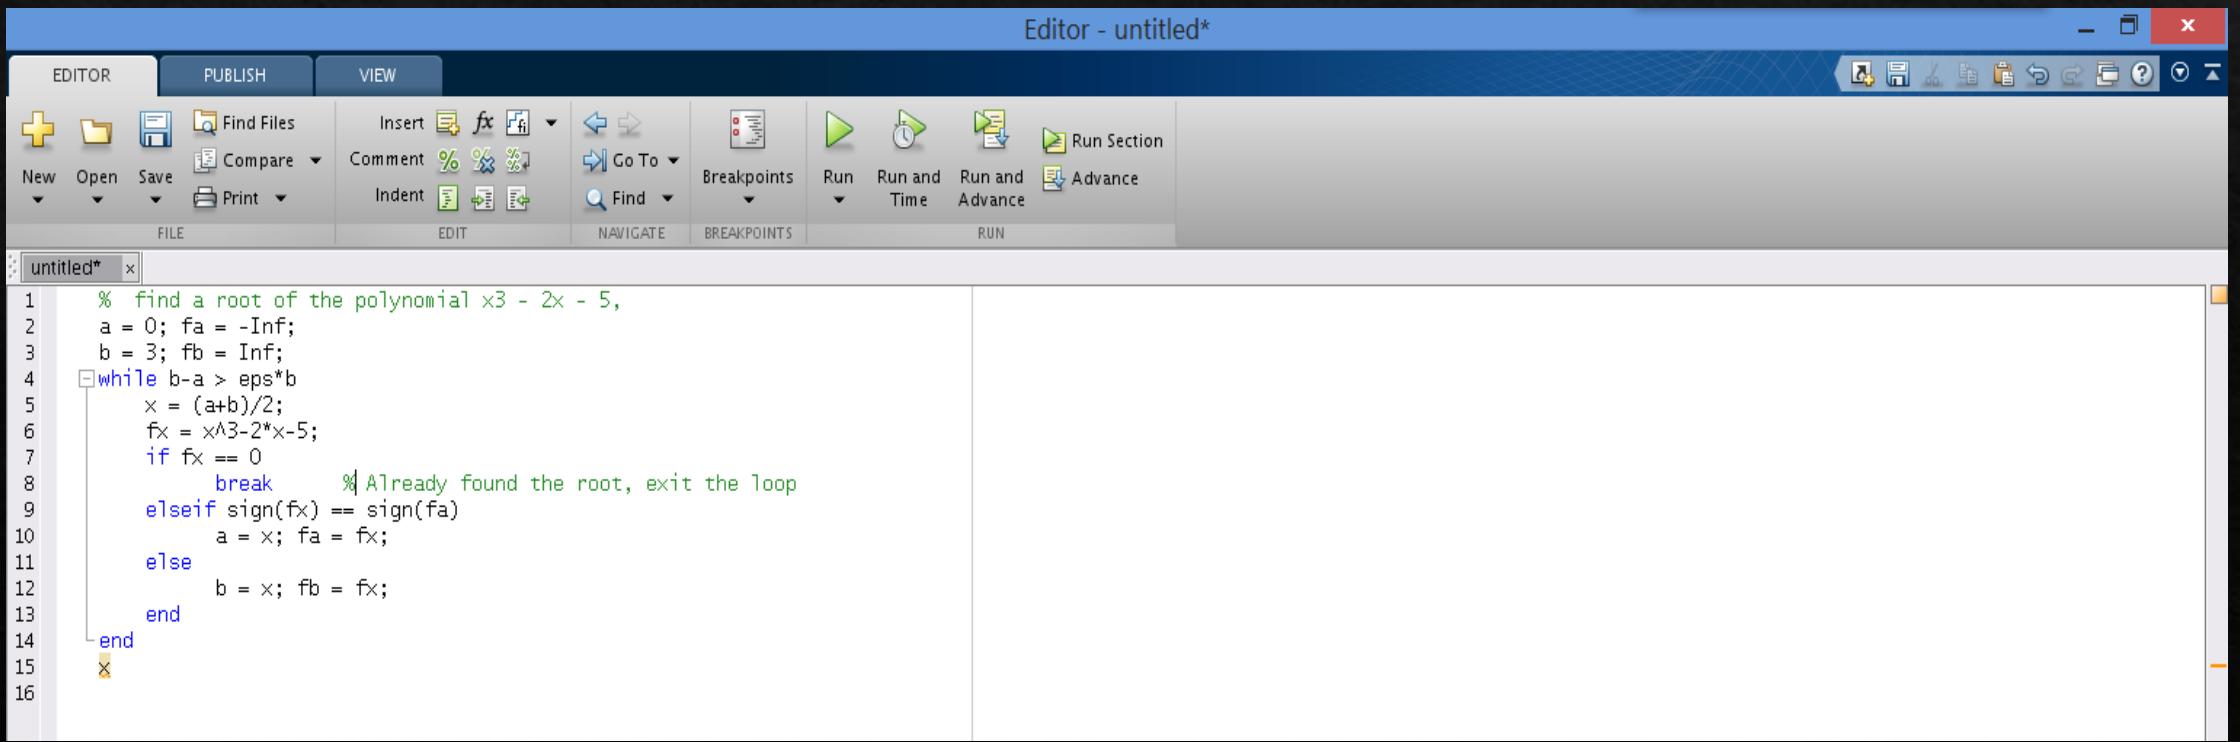


- Use VNC to speed up graphical interface.

Refer to: <https://www.bu.edu/tech/support/research/system-usage/getting-started/remote-desktop-vnc/>

M-file

- ❖ An m-file is a simple text file where you can place MATLAB commands.
- ❖ Save your works
- ❖ Convenient for debugging
- ❖ Run directly. Pre-compiling is unnecessary.



The screenshot shows the MATLAB Editor window titled "Editor - untitled*". The menu bar includes "EDITOR", "PUBLISH", and "VIEW". The toolbar contains icons for New, Open, Save, Find Files, Compare, Comment, Indent, Go To, Breakpoints, Run, Run and Time, Run and Advance, and Advance. The code editor window displays the following MATLAB script:

```
1 % find a root of the polynomial x^3 - 2x - 5,
2 a = 0; fa = -Inf;
3 b = 3; fb = Inf;
4 while b-a > eps*b
5     x = (a+b)/2;
6     fx = x^3-2*x-5;
7     if fx == 0
8         break      % Already found the root, exit the loop
9     elseif sign(fx) == sign(fa)
10        a = x; fa = fx;
11    else
12        b = x; fb = fx;
13    end
14 end
```

Text platform

```
$ matlab -nodisplay          % Work in text interface. Does not display any graph.  
$ matlab -nodesktop           % Program in text interface. Pop out graphs when necessary.
```

- ❖ Many Linux commands (some prefixed with an exclamation mark) are available within Matlab platform, such as:

cd, ls, pwd, !cp, !rm, !mv, !cat, !vim, !diff, and !grep

- ❖ Edit M-file and run the program:

```
>> !vim mfilename.m           % edit in text window  
>> edit mfilename.m          % create a new or open an existing m-file in graphical window  
>> open mfilename.m          % open an existing m-file in graphical window  
>> run mfilename.m           % run the program  
>> mfilename                 % run the program
```

Submit a batch job to background

- ❖ Submit a batch job to background using a script

```
$ qsub script.sh
```

- ✓ Write a batch job script using any Linux editor (such as vim, emacs, gedit, or nano).
- ✓ A typical batch script for a Matlab job (Supposed that source codes are saved in an m-file):

```
#!/bin/bash -l      # Start a bash script. The option -l is necessary to enable module tool.  
#$ -pe omp 1       # Request 1 CPU core  
#$ -l h_rt=12:00:00 # Request wall time  
#$ -N jobname       # Give a job name  
module load matlab   # Set up environment variables for running matlab  
matlab -nodisplay -singleCompThread -r "addpath /path/to/work; mfile_name"  # Run program
```

Exercise 1

- ❖ Run a Matlab program on BU SCC:

Request for an interactive session first.

- i) Write a Matlab code in an m-file to print “Hello world” (e.g. use the **disp** function).
- ii) Run the program interactively in Matlab platform.
- iii) Submit a batch job to run the program in the background.

Standalone executable

- ❖ Create a standalone executable

```
$ mcc -mv -o myexe name.m
```

- ✓ **-mv** produces a standalone and shows actions taken
- ✓ **-o myexe** specifies the executable name
- ✓ The **name** of the m file must be the same as the main function name in it.
- ✓ A script named **run_myexe.sh** is automatically created for running jobs in Linux bash shell.

- ❖ Run the executable (in Linux shell environment)

```
$ module load mcr    # Set up environment variables for running the executable.
```

```
$ mcr ./myexe    # Execute the executable using MATLAB Compiler Runtime (MCR).
```

- ✓ There is no considerable performance difference between running a standalone executable and running an m-file directly. The later one is more recommended since it is simpler.

Distributed jobs

- ❖ Submit multiple jobs using one command line.
- ❖ Distribute independent jobs:

```
$ qsub -t 1-4 script.sh
```

- ✓ The batch system sets up the SGE_TASK_ID environment variable which can be used to pass the task ID to the program, for example:

```
matlab -nodisplay -singleCompThread -r "rand($SGE_TASK_ID); exit"
```

```
matlab -nodisplay -singleCompThread -r "id='\$SGE_TASK_ID'; disp(id); exit"
```

```
matlab -nodisplay -singleCompThread -r "id=getenv('SGE_TASK_ID'); disp(id); exit"
```

- ❖ Distribute dependent jobs:

```
$ qsub -N job1 script1.sh
```

```
$ qsub -N job2 -hold_jid job1 script2.sh
```

- ✓ Job 2 does not start until job1 ends.

Optimize Matlab codes

- ❖ Tools for measuring performance and optimizing codes
- ❖ Remove unnecessary work
- ❖ Optimize memory usage
- ❖ Use built-in functions/operators

Tools for measuring performance and optimizing codes

```
>> tic    % Start measuring time  
>> toc    % End measuring time  
>> timeit (function_name)    % Measure time required to run function  
>> mlint ('mfile_name')    % Reports potential problems and opportunities for code optimization  
>> profile on    % Turn on profile (before the program starts).  
>> profile viewer    % View the results in the Profiler window (after the program ends).
```

Unnecessary work (1): redundant operations

- ❖ Avoid redundant operations in loops

```
for i=1:N  
    x = 10;  
    .  
    .  
end
```

bad

```
x = 10;  
for i=1:N  
    .  
    .  
end
```

good

Unnecessary work (2): reduce overhead

..from function calls

bad

```
function myfunc(i)
    % do stuff
end

for i=1:N
    myfunc(i);
end
```

good

```
function myfunc2(N)
    for i=1:N
        % do stuff
    end
end

myfunc2(N);
```

..from loops

bad

```
for i=1:N
    x(i) = i;
end
for i=1:N
    y(i) = rand();
end
```

good

```
for i=1:N
    x(i) = i;
    y(i) = rand();
end
```

Unnecessary work (3): logical tests

Avoid unnecessary logical tests...

...by using short-circuit
logical operators

```
if (i == 1 | j == 2) & k == 5  
    % do something  
end
```

bad

```
if (i == 1 || j == 2) && k == 5  
    % do something  
end
```

good

...by moving known cases
out of loops

bad

```
for i=1:N  
    if i == 1  
        % i=1 case  
    else  
        % i>1 case  
    end  
end
```

good

```
% i=1 case  
for i=2:N  
    % i>1 case  
end
```

Unnecessary work (4): reorganize equations

Reorganize equations to use fewer or more efficient operators

Basic operators have different speeds:

Add	3- 6 cycles
Multiply	4- 8 cycles
Divide	32-45 cycles
Power, etc	(worse)

bad

```
c = 4;  
for i=1:N  
    x(i)=y(i)/c;  
    v(i) = x(i) + x(i)^2 + x(i)^3;  
    z(i) = log(x(i)) * log(y(i));  
end
```

good

```
s = 1/4;  
for i=1:N  
    x(i) = y(i)*s;  
    v(i) = x(i)*(1+x(i)*(1+x(i)));  
    z(i) = log(x(i) + y(i));  
end
```

Unnecessary work (5): avoid re-interpreting code

MATLAB improves performance by interpreting a program only once, unless you tell it to forget that work by running “clear all”

Value of ItemType	Items Cleared							
	Variables in scope	Scripts and functions	Class definitions	Persistent variables	MEX functions	Global variables	Import list	Java classes on the dynamic path
all	✓	✓		✓	✓	✓	From command prompt only	
variables	✓							

Memory (1): preallocate arrays

- ◆ Arrays are always allocated in contiguous address space.
- ◆ If an array changes size, and runs out of contiguous space, it must be moved. For example,

```
x = 1;  
for i = 2:4  
    x(i) = i;  
end
```

- ◆ This can be very very bad for performance when variables become large.

Memory Address	Array Element
1	x(1)
...	...
2000	x(1)
2001	x(2)
2002	x(1)
2003	x(2)
2004	x(3)
...	...
10004	x(1)
10005	x(2)
10006	x(3)
10007	x(4)

Memory (1): preallocate arrays

- ❖ Preallocating array to its maximum size prevents intermediate array movement and copying.

```
A = zeros(n,m); % initialize A to 0  
A(n,m) = 0;       % or touch largest element
```

- ❖ If maximum size is not known, estimate with upper bound. Remove unused memory after.

```
A=rand(100,100);  
% . . .  
% if final size is 60x40, remove unused portion  
A(61:end,:)=[]; A(:,41:end)=[]; % delete
```

Memory (2): for-loop order

- ❖ It is faster to access continuous memory addresses than separated ones.
- ❖ Multidimensional arrays are stored in memory along columns (column-major).
- ❖ Best if inner-most loop is for array left-most index, etc.

bad

```
n=5000; x = zeros(n);
for i = 1:n      % rows
    for j = 1:n    % columns
        x(i,j) = i+(j-1)*n;
    end
end
```

good

```
n=5000; x = zeros(n);
for j = 1:n      % columns
    for i = 1:n    % rows
        x(i,j) = i+(j-1)*n;
    end
end
```

Memory (3): avoid unnecessary variables

- ❖ Avoid time needed to allocate and write data to main memory.
- ❖ Compute and save array in-place improves performance and reduces memory usage.

bad

```
x = rand(5000);  
y = x.^2;
```

good

```
x = rand(5000);  
x = x.^2;
```

Exercise 2

- ❖ Optimize this Matlab code to obtain a better performance.
- ✓ Hint: Use **mlint** to report potential problems and opportunities for code optimization.

```
n=5000;  
for i=1:n  
    for j=1:n  
        a=3./2.;  
        if (i==1)  
            x(i, j) = 5.;  
        else  
            x(i, j) = i * log(2.) * log(a) + j^2 / 2.;  
        end  
    end  
end  
y=x.^2;
```

Matlab built-in functions/operators

- ❖ Some useful Matlab built-in functions/operators:

Matrix operations: *, mtimes, inv, eig

Solve linear equation: mldivide, linsolve, \

Decomposition: lu, qr

Optimization: fminsearch, fzero

- ✓ A full list:

<https://www.mathworks.com/help/matlab/functionlist.html#linear-algebra>

- ❖ Built-in functions/operators are optimized and performs well in general.

Matrix multiplication

- ❖ The built-in operator `*` has been optimized and has much better performance.

```
n = 500
for j=1:1:n      % initialize data
    for i=1:1:n
        A(i,j)=i+j;  B(i,j)=2*i-j;
    end
end
C=zeros(n); D=zeros(n);
tic
for i=1:1:n      % mattix multiplication by loops
    for j=1:1:n
        C(i,j)= A(i,:)*B(:,j);
    end
end
toc
tic
D = A * B;    % mattix multiplication by built-in operator
toc
isequal(C,D)   % Check results
```

The backslash operator: solve linear equations

- ❖ Solve the linear system $A^*x = b$

$x = A \backslash b$

$x = \text{mldivide}(A, b)$

$x = \text{linsolve}(A, b)$

$x = \text{linsolve}(A, b, \text{opts})$

$A = \text{triu}(\text{rand}(5, 5));$ % random 5*5 up-triangle matrix

$b = \text{rand}(5, 1);$ % random column array

$\text{opts.UT} = \text{true};$ % up-triangle is true

$x = \text{linsolve}(A, b, \text{opts})$

$x = A \backslash b$ % What does the backslash operator actually do behind the scene?

Decomposition

❖ LU decomposition

$[L,U] = \text{lu}(A)$ % expresses a matrix A as the product of two essentially triangular matrices, one of them a permutation of a lower triangular matrix and the other an upper triangular matrix. The decomposition is often called the LU, or sometimes the LR, decompostion.

$[L,U] = \text{lu}(A);$ % obtain L and U matrices

$y=L\backslash b;$

$x=U\backslash y$ % These 3 lines together are equivalent to $x=A\backslash b$

❖ QR decomposition

$[Q,R] = \text{qr}(A)$ % expresses a matrix A into a product $A = QR$ of an orthogonal matrix Q and an upper triangular matrix R.

$[Q,R] = \text{qr}(A);$ % obtain Q and R matrices

$y=Q'^{*}b;$

$x=R\backslash y$

Exercises 3

- ❖ **Exercise 3.1:** Create an $n \times n$ symmetric matrix A in one of the following ways:
 - i) use built-in function *triu* (up-triangle matrix) and the matrix transpose operation '`;`';
 - ii) use control flow (*for*, *if*, *else*, ...).
- ❖ **Exercise 3.2:** Given a symmetric matrix A from exercise 3.1, solve the linear algebra equation $A^*x=b$ using the following two methods:
 - i) use the matrix inverse function *inv* .
 - ii) use the backslash operator `\` .

Compare the computational time and numerical error of the two cases.

Further Information

- ❖ MathWorks Web:
 - ✓ MATLAB documentation: <http://www.mathworks.com/help/matlab/>
- ❖ BU Research Computing Services (RCS) Web:
 - ✓ MATLAB basics: <http://www.bu.edu/tech/support/research/software-and-programming/common-languages/matlab/>
- ❖ A book: *Accelerating MATLAB Performance: 1001 tips to speed up MATLAB programs* by *Yair M. Altman* .
- ❖ RCS help: help@scc.bu.edu , shaohao@bu.edu